

ADL HW3 Report

Q1: Retriever & Reranker Tuning (5%)

Retriever Training Process (2.5%)

1. Training data

The retriever model is trained using query–passage pairs derived from a labeled dataset.

Anchor: Each query.

Positive sampling: For each query, we retrieve its relevant passages as positive sample.

Negative sampling: Instead of explicit negatives, we use the *MultipleNegativesRankingLoss* to automatically treats other positive passages within the same batch as negative examples for the current query.

2. Loss function

I use *MultipleNegativesRankingLoss* as the contrastive loss function.

The loss for query i is:

$$L_i = -\log \frac{\exp(\text{sim}(q_i, p_i)/\tau)}{\sum_j \exp(\text{sim}(q_i, p_j)/\tau)}$$

where:

$\text{sim}(\cdot, \cdot)$ is cosine similarity,

p_i is the positive passage,


$p_j, (j \neq i)$ are in-batch negatives,

τ is a temperature hyperparameter (implicitly handled inside the implementation).

3. Hyperparameters

Hyperparameter	Value
Base model	intfloat/multilingual-e5-small
Batch size	32
Epochs	1
Learning rate	2e-5 (default)
Warmup steps	10% of total steps
Evaluation steps	Every 100 steps
Loss function	MultipleNegativesRankingLoss
Optimizer	Default
Seed	42
Mixed precision	Enabled (use_amp=True)

4. Training curve



The retriever achieves recall@10 = 0.8687 on evaluation set after 1 epoch.
(0.8256 on public test set.)

Reranker Training Process (2.5%)

1. Training Data

The reranker is trained using query-passage pairs in a supervised binary classification setting.

Each training sample is created from three components:

Anchor: a query q (from the training queries).

Positive sample: one relevant passage p^+ , determined by non-zero relevance labels in the qrels file.

Negative samples: one (or more) irrelevant passages p^- , randomly sampled from the corpus while excluding all known relevant passages.

2. Loss Function

The reranker uses the Cross-Entropy Loss implemented internally in the Sentence Transformers CrossEncoder class.

Given the query-passage pair, the model outputs a single relevance score in $[0, 1]$, and the loss encourages correct binary classification:

$$\mathcal{L} = -[y \log(\hat{y}) + (1 - y) \log(1 - \hat{y})]$$

where:


$y \in \{0, 1\}$ is the true label (relevant / non-relevant),

\hat{y} is the model's predicted probability.

3. Hyperparameters

Parameter	Value
Base model	cross-encoder/ms-marco-MiniLM-L-12-v2
Lose function	Cross-Entropy Loss
train_epochs	1
train_batch_size	8
learning_rate	2e-5
warmup_steps	100
num_neg_per_pos	1
max_len	512
eval_steps	1000
top_k	10
eval_batch_q	16
seed	42

4. Training curve



MRR@10 rises from 0.6275 (without reranker) to 0.7375 (with reranker) on evaluation set. (0.7302 on public test set.)

Q2: Prompt Optimization (3%)

Provide a detailed explanation of how you designed your prompt. (1.5%)

I tried three methods: basic, CoT, confidence-base.

Method 1: *"Answer the user concisely..."* enforces short, direct responses.

Method 2: *"Explain your reasoning briefly..."* encourages chain-of-thought.

Method 3: *"Provide an answer only if the context clearly supports it; otherwise, respond with 'CANNOTANSWER'."* tells model how to handle situations when there's no answer.


Present at least three different prompts you experimented with. (1.5%)

Method	Prompt	CosSim
1 Basic	<pre>def get_inference_system_prompt(): return "Answer the user concisely based on the context passages." def get_inference_user_prompt(query, context_list): return f"Question: {query}\n\nContext:\n" + "\n\n".join(context_list)</pre>	0.3660
2 CoT	<pre>def get_inference_system_prompt2(): return "You are an assistant that answers questions using evidence from the given passages. Explain your reasoning briefly before giving the final answer." def get_inference_user_prompt2(query, context_list): return f"Question: {query}\n\nRelevant Passages:\n" + "\n---\n".join(context_list) + "\n\nPlease reason step by step and conclude with a short final answer."</pre>	0.3397
3 confidence base	<pre>def get_inference_system_prompt3(): return "Provide an answer only if the context clearly supports it; otherwise, respond with 'CANNOTANSWER'." def get_inference_user_prompt3(query, context_list): return f"Question: {query}\n\nSupporting Contexts:\n" + "\n\n".join(context_list)</pre>	0.3572

Basic method is the best; CoT takes long time when inferencing; Teach model “CANNOTANSWER” does not help much.

Q3: Additional Analysis (2%)

I analyzed the distribution of CosSim on all public test data. Furthermore, I separate them into hit@3 and miss@3, in the other words, correct paragraph was sent to LLM or not. Unsurprisingly, CosSim of hit@3 is higher then miss@3. Note that no one is below -0.1.




	hit@3	miss@3	total
count	2658	684	3342
mean	0.4000	0.2046	0.3600
std	0.2455	0.1747	0.2458

Bonus: RL in the loop (2%)

I adopted the A2C (Advantage Actor-Critic) algorithm from stable-baselines3, and trained the agent in a custom environment where the reward was based on the Cosine Similarity between the generated and gold answers.

The learned action distribution showed that the model preferred larger context sizes, $\text{top_m} = 6$ appears most times.

If top_m is too low, the correct paragraphs may not be sent to LLM. If top_m is too high, too many irrelevant paragraphs will be included. Dynamically adjusting the number of passages helps the LLM achieved higher performance. (CosSim 0.3660 → 0.3910)



Note

All trainings are on workstation with $1 \times$ RTX A6000 GPU.

References

1. ChatGPT-5
2. Gemini 2.5 Pro
3. Wang, L., Yang, N., Wang, X., Joty, S., & Lin, J. (2022). Text Embeddings by Reranking. arXiv preprint arXiv:2212.03534.
4. Reimers, Nils, and Iryna Gurevych. "Sentence-bert: Sentence embeddings using siamese bert-networks." arXiv preprint arXiv:1908.10084 (2019).
5. Karpukhin, Vladimir, et al. "Dense Passage Retrieval for Open-Domain Question Answering." EMNLP (1). 2020.
6. Mnih, Volodymyr, et al. "Asynchronous methods for deep reinforcement learning." International conference on machine learning. PMLR, 2016.